



THE PROBLEM

Consulting's Biggest Bottleneck Is Manual Work – and Lost Institutional Knowledge

Analysts at top consulting firms spend **60–70% of their time** on repetitive, low-value tasks: research synthesis, storyline structuring, and deck formatting. But the deeper bottleneck is that they also can't leverage their firm's institutional memory – past presentations, client feedback, and proven frameworks sit locked in drives instead of accelerating new work. The result is slow, inconsistent output that doesn't scale.

ConsultFlow solves both problems by reducing manual work and unlocking organizational knowledge.

ConsultFlow AI: An End-to-End Platform for Consulting Decks

ConsultFlow AI automates the entire deck lifecycle — from client discovery through research, QA, storyline generation, and final delivery in Google Slides. It is **RAG-powered** and **framework-aware from day one**, grounded in organizational knowledge and built to apply consulting best practices automatically. These are not future features — they are live capabilities for consulting teams who need speed without sacrificing quality.



RAG-Powered Research

Automated synthesis grounded in organizational knowledge and historical presentations



Pre-Built Frameworks

Pyramid Principle, MECE, and Porter's Five Forces automatically applied



Human-in-the-Loop QA

Analyst approval gates at every critical stage




Instant Deck Delivery

Auto-rendered Google Slides delivered to Drive

The 6-Gate Pipeline: Linear, Transparent, Human-Controlled

Every engagement flows through a structured six-stage pipeline. Analysts retain full control with one-click approval at each gate — AI does the heavy lifting, humans make the decisions.



Revision notes at any stage are fed back into the AI for intelligent regeneration — creating a continuous improvement loop without restarting the workflow.

Modern, Serverless, Production-Ready Stack

RAG is live now and sits at the center of the system — grounding every AI output in organizational knowledge before generation happens.

RAG Layer — Live Now

Retrieval-Augmented Generation grounds every AI output in organizational knowledge. Historical presentations, client research, past engagement insights, and firm-specific frameworks are automatically retrieved and injected into the AI context to ensure consistency, reduce hallucinations, and leverage institutional memory.

Sources: Drive-based historical presentations, past client feedback, organizational research libraries, and pre-loaded consultant frameworks like Pyramid Principle, MECE, and Porter's Five Forces.

Status: LIVE NOW — this retrieval layer is already active in production.

Core Stack

Frontend & Hosting: TanStack Start on Cloudflare Workers — edge-native, sub-100ms response times, zero cold starts.

Backend Automation: n8n self-hosted on Railway orchestrates all AI agent workflows with full observability.

AI Model: Gemini 2.5 Flash — fast, cost-efficient, and capable of long-context reasoning for complex research synthesis.

Database & Auth: Supabase (PostgreSQL + Row Level Security) provides secure, multi-tenant data isolation.

Deck Output: Google Slides API with QuickChart for auto-generated, formatted data visualizations.

One Workspace. Every Client Engagement.

The ConsultFlow dashboard is the command center for consulting teams. Every engagement is tracked in real time as it moves through the pipeline — from initial discovery to final deck delivery.

Status Tracking


Engagements display live status: Discovery → Research → QA → Storyline → Delivered. No guessing where work stands.

One-Click Approval Gates

Analysts approve or request revisions at each stage with a single click — no context switching, no friction.

Revision Feedback Loop

Revision notes are fed directly back into the AI for intelligent regeneration, preserving context and intent.



Research That Reads Like a Consultant Wrote It

The AI Research Agent is RAG-powered and grounded in organizational frameworks, so it doesn't produce a wall of text. It synthesizes market context, key findings with source citations, strategic recommendations, and risk factors into a proper consulting research brief.

→ Market Context

Industry landscape, competitive dynamics, and macro trends — enriched by live market data and external sources.

→ Key Findings

Evidence-backed insights with inline source citations, informed by historical presentations and past client research from your Drive.

→ Strategic Recommendations

Actionable next steps aligned to client objectives, structured by pre-loaded consultant frameworks like the Pyramid Principle and MECE.


→ Risk Register

Identified risks with mitigation framing, using organizational knowledge and frameworks to keep analysis consistent with past work.

Every research output is grounded in three layers: live market data and external sources, historical presentations and past client research from your Drive, and pre-loaded consultant frameworks such as the Pyramid Principle, MECE, and Porter's Five Forces. This ensures consistency with past work and automatically leverages institutional knowledge.

From Approved Storyline to Formatted Deck in Seconds

Once the storyline is approved, ConsultFlow automatically renders a fully formatted Google Slides deck — delivered directly to the analyst's Drive. No manual copy-paste. No formatting cleanup.



The SCR (Situation-Complication-Resolution) framework is the backbone of every storyline — the same narrative structure used by top-tier consulting firms worldwide. The AI generates slide-by-slide content with clear headers, supporting data points, and logical flow.

Eval System: Scoring Every AI Output Automatically

In agentic systems, you can't manually review every output. ConsultFlow's eval system scores every AI-generated artifact across three dimensions before it reaches the analyst.

3

Evaluation Dimensions

Completeness, Groundedness, and Overall
Score assessed on every output

7/10

Pass Threshold

Outputs below this score are flagged and
regenerated before reaching the analyst

100%

Auto-Scored

Every AI output is evaluated automatically —
zero manual review required

Completeness

Did the AI cover all required sections defined in the discovery brief?
Missing sections are flagged immediately.

Groundedness

Are all claims backed by cited sources and contextual evidence?
Hallucinated or unsupported assertions are caught.

Evals Are What Separate a Demo from Production



Without evals, agentic AI systems are black boxes. You don't know when quality degrades, which stages underperform, or whether improvements are working. ConsultFlow's eval system solves all three.

1

Consistent Quality

Confidence that AI output meets a defined standard across every engagement

2

Stage Visibility

Pinpoint exactly which pipeline stages underperform and need prompt tuning

3

Continuous Improvement

A structured feedback loop to iteratively improve prompts and reduce rework

ROADMAP

The Foundation Is Built – Now It Scales

ConsultFlow's core pipeline is production-ready. The next phase focuses on enterprise hardening, deeper personalization, and team collaboration features.



Google OAuth & Per-User Drive Access

Individual Drive integration so every analyst works in their own environment with proper permissions.



Supabase Edge Functions

Replacing n8n with native Edge Functions for lower latency and tighter Supabase integration.



RAG Pipeline for Client Knowledge

Client-specific knowledge bases that ground AI outputs in proprietary research and past engagements.



Multi-User Access & Role Management

Team workspaces with analyst, reviewer, and admin roles — built for firm-wide deployment.